

GO-Bayes: Gene Ontology-based overrepresentation analysis using a Bayesian approach

Song Zhang^{1,*}, Jing Cao², Y. Megan Kong³ and Richard H. Scheuermann^{1,3,*}

¹Department of Clinical Sciences, U.T. Southwestern Medical Center, 5323 Harry Hines Boulevard Dallas, TX 75390-9072, ²Department of Statistical Science, Southern Methodist University, Dallas, TX 75275-0332,

³Department of Pathology, U.T. Southwestern Medical Center, 5323 Harry Hines Boulevard Dallas, TX 75390-9072, USA

Associate Editor: David Rocke

ABSTRACT

Motivation: A typical approach for the interpretation of high-throughput experiments, such as gene expression microarrays, is to produce groups of genes based on certain criteria (e.g. genes that are differentially expressed). To gain more mechanistic insights into the underlying biology, overrepresentation analysis (ORA) is often conducted to investigate whether gene sets associated with particular biological functions, for example, as represented by Gene Ontology (GO) annotations, are statistically overrepresented in the identified gene groups. However, the standard ORA, which is based on the hypergeometric test, analyzes each GO term in isolation and does not take into account the dependence structure of the GO-term hierarchy.

Results: We have developed a Bayesian approach (GO-Bayes) to measure overrepresentation of GO terms that incorporates the GO dependence structure by taking into account evidence not only from individual GO terms, but also from their related terms (i.e. parents, children, siblings, etc.). The Bayesian framework borrows information across related GO terms to strengthen the detection of overrepresentation signals. As a result, this method tends to identify sets of closely related GO terms rather than individual isolated GO terms. The advantage of the GO-Bayes approach is demonstrated with a simulation study and an application example.

Contact: song.zhang@utsouthwestern.edu;
richard.scheuermann@utsouthwestern.edu

Supplementary information: Supplementary data are available at *Bioinformatics* online.

Received on August 4, 2009; revised on December 21, 2009; accepted on February 3, 2010

1 INTRODUCTION

In typical high-throughput experiments, such as gene expression microarrays, the first step in the analysis of the results is often to produce groups of genes based on certain criteria (e.g. genes that are differentially expressed). To gain more mechanistic insights into the underlying biology, overrepresentation analysis (ORA) is conducted to use the knowledge of the functional characteristics of the genes to investigate whether gene sets associated with particular biological functions are overrepresented in the identified gene groups. Drăghici *et al.* (2003) is the first paper to discuss the overrepresentation

problem and propose different statistical methods that can be used in this area. ORA is based on the postulate that if a biological process has more identified genes than expected by chance alone, that biological process is probably linked to the experiment.

One of the most popular gene description databases used in ORA was developed by the Gene Ontology (GO) Consortium (Ashburner *et al.*, 2000). Each GO term annotates a set of genes, indicating their known molecular functions, involvement in biological processes and cellular component locations. GO terms are structured in a directed acyclic graph (DAG) of parent–child relationship, where a child indicates a more specific biological classification than its parent(s). Based on the *true-path* rule the annotation of a gene to a GO term implies automatic annotation to all the ancestors of that term. Another feature is that a GO term is allowed to have more than one parent nodes, a feature known as multiple inheritance. For example, immune response is not only a specific form of organismal movement but also a part of defense response. Furthermore, some genes might be annotated by a parent GO node but not by any of its children because less is known about that gene's specific function. For a rigorous analysis, these dependency characteristics of the GO DAG need to be considered when developing statistical methods to detect overrepresentation of GO annotations.

In ORA, the most commonly used statistical test is based on the hypergeometric distribution or its binomial approximation (Cho *et al.*, 2001; Khatri *et al.*, 2002; Drăghici *et al.*, 2003; Al-Shahrour *et al.*, 2004; Beissbarth and Speed, 2004; Lee *et al.*, 2005; Lee *et al.*, 2006; Luo *et al.*, 2007; among others). Let A denote a GO term or the set of genes annotated to A (with cardinality I_A), and let S denote the set of genes (with cardinality I_S) based on a certain criterion (i.e. differential expression) from a full gene list G (with cardinality I) in an experiment. The number of genes belonging to both S and A ($S \cap A$), denoted by n_A , indicates the representation of A in S . Under the null hypothesis that S and A are independent (i.e. the GO term is irrelevant to the gene cluster), n_A follows a hypergeometric distribution. The P -value measuring the significance of association is the tail probability of observing n_A or more genes annotated by A in S ,

$$P\text{-value} = \sum_{k=n_A}^{\min(I_A, I_S)} \frac{\binom{I_A}{k} \binom{I-I_A}{I_S-k}}{\binom{I}{I_S}}, \quad (1)$$

where $\binom{m}{n} = \frac{m!}{n!(m-n)!}$ is the binomial coefficient. Many software and webtools (Onto-Express, CLASSIFI, GoMiner, EASEonline,

*To whom correspondence should be addressed.

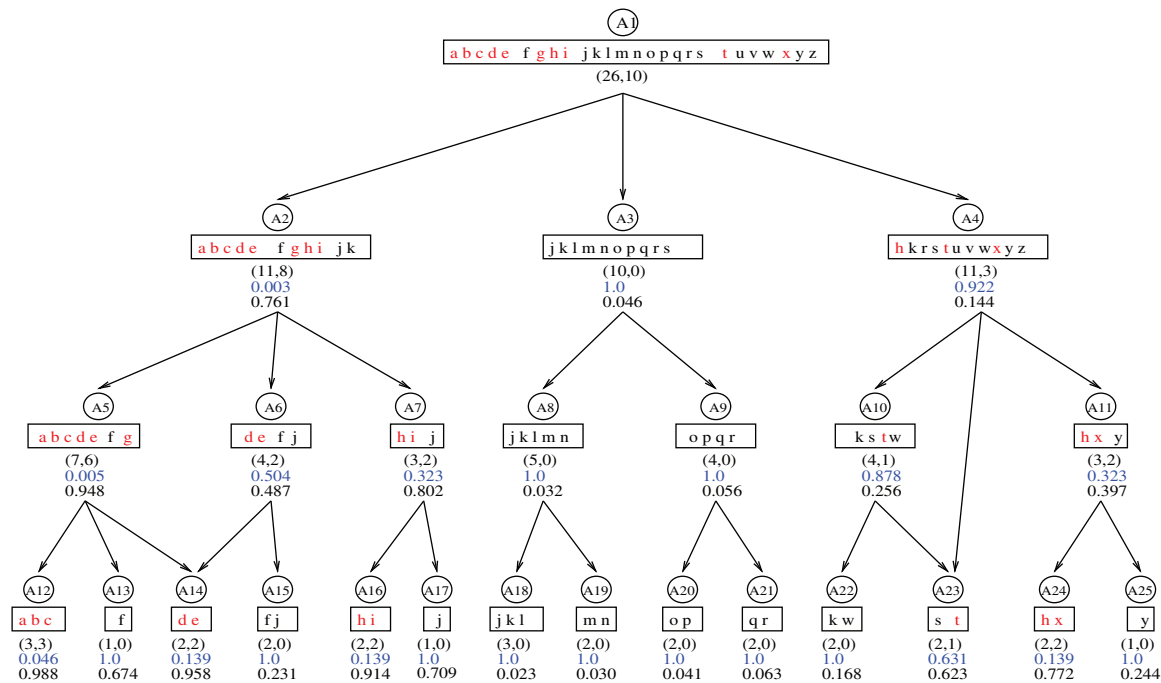


Fig. 1. Comparison of the P -value and the B -score with an artificial DAG. The full list of genes (G) are denoted as lowercase letters; the genes in set S are marked in red. The rectangles contain the subset of genes annotated by each node, where (I_A, n_A) are listed under each rectangle. The hypergeometric P -values are listed in blue and the B -scores are listed in black underneath the P -values.

GeneMerge, FuncAssociate, GOTree Machine, etc.) have been developed based on the hypergeometric P -value. Detailed review can be found in Khatri and Drăghici (2005).

The hypergeometric P -value provides a straightforward measure of overrepresentation for each individual GO term. However, the major drawback of this approach is that it ignores the hierarchical structure in the GO DAG, which contains a substantial amount of information regarding the interactions among the GO terms.

We use an artificial DAG (Fig. 1) to illustrate this issue. It consists of 25 nodes $\{A_j, j=1, \dots, 25\}$, each denoted by a circle. Let $G = \{a, b, \dots, z\}$ denote the full list of genes, and 10 genes in S are marked in red. The rectangles contain the subset of genes annotated by each node, where (I_A, n_A) are listed under each rectangle. Figure 1 includes some important features of the GO DAG, such as multiple inheritance (e.g. A_{14} and A_{23} have two parent nodes) and that a gene might be annotated at different specific level (eg. gene k is annotated by A_2 , but not by any of its children: A_5 , A_6 and A_7). Note that in this example S is overrepresented in the regions under A_2 , while underrepresented in the regions under A_3 and A_4 . We report the hypergeometric P -values estimated by (1) in blue.

First, given I and I_S , GO terms with the same (I_A, n_A) have exactly the same hypergeometric P -values. For example, with $(I_A, n_A) = (2, 2)$, the P -values for A_{14} and A_{24} are both 0.139. Based on Figure 1, A_{14} is more likely to be linked with S because of the stronger evidence of overrepresentation in its neighboring nodes (related biological functions). Using evidence only from individual terms, the hypergeometric P -value does not differentiate between A_{14} and A_{24} . Second, given I and I_S , the hypergeometric P -value has a lower limit determined by I_A , which is denoted

as $L(I_A)$. Specifically, $L(I_A)$ is achieved when $n_A = I_A$, i.e. the P -value reaches its lower limit when all the genes annotated by A are in S . The smaller the I_A , the larger the $L(I_A)$. For example, we have $L(3) = 0.046$ as in A_{12} and $L(2) = 0.139$ as in A_{14} . This observation suggests that if we set the threshold for the P -value at $L(k)$, then the hypergeometric test could not identify any GO terms with $I_A < k$. In ORA, detecting more specific GO terms, which usually have a relatively smaller I_A , might be more desirable because they provide more detailed biological information. However, the hypergeometric test tends to identify less specific GO terms because of the constraint of $L(I_A)$. In Figure 1, the most significant term selected by the P -value is A_2 , a term next to the root. The more specific terms that are associated with S (i.e. A_{12} and A_{14}) are considered less significant compared with A_2 . All of these issues essentially stem from the limitation of the hypergeometric test in treating the GO terms as independent entities and ignoring their interrelated structure.

Recently some new methods have been proposed in ORA. Lewin and Grieve (2006) propose to group closely related GO nodes together and then compute a hypergeometric P -value for each group. In their approach, the graphical distance between nodes in the GO DAG is assumed to have some quantitative biological meaning. Alexa *et al.* (2006) calculate the overrepresentation of a GO term from leaf to root by downweighting the contribution of genes which are annotated by a child term that has been found to be significantly enriched. Grossmann *et al.* (2007) measure the overrepresentation of each GO term relative to its parent(s). From root to leaf, their method computes the significance of overrepresentation for a GO term conditional on the representation at the parent term(s). The last

two methods purposely remove the dependence between parent and child terms.

In this article, we develop a Bayesian hierarchical model to incorporate the dependence structure of the DAG in assessing GO term overrepresentation. It takes into account evidence not only from individual GO terms, but also from their related terms (i.e. parents, children, siblings, etc.). The Bayesian framework enables borrowing information across related GO terms to strengthen the detection of overrepresentation signals. As a result, this method tends to identify sets of closely related GO terms rather than individual unrelated GO terms. The utility of the method is demonstrated using a gene expression microarray dataset from a human B cell stimulation experiment.

2 METHOD

2.1 The Bayesian model

The proposed method is called GO-Bayes: GO-based ORA using a Bayesian approach. In the model, each GO term has a relevance parameter measuring its association with the selected genes in S . The novelty of the model is that the complex dependence structure in the GO DAG is incorporated via a hierarchical prior on the relevance parameters.

To introduce GO-Bayes, we define the following notations. Let $\mathbf{A} = \{A_j, j=1, \dots, J\}$ be the set of GO terms involved in the annotation of the full list of I genes (i.e. $I_{A_j} > 0$, for $j=1, \dots, J$), among which I_S genes are grouped in S in the experiment. We use $A_{j_1} \rightarrow A_{j_2}$ to indicate that A_{j_1} is a parent of A_{j_2} . We define $\mathbf{P}_j = \{A_k : A_k \rightarrow A_j\}$ and $\mathbf{C}_j = \{A_k : A_j \rightarrow A_k\}$ to be the sets of parent and child nodes of A_j , respectively. We use $|\mathbf{U}|$ to denote the cardinality of set \mathbf{U} . Without loss of generality, let A_1 denotes the root node and $|\mathbf{P}_1| = 0$. For the non-root nodes, we have $|\mathbf{P}_j| \geq 1$, where $|\mathbf{P}_j| > 1$ indicates multiple inheritance. We call A_j an end node if $|\mathbf{C}_j| = 0$ and an inner node otherwise. For example, in Figure 1, A_1 is the root node, $\mathbf{P}_{10} = \{A_4\}$, $\mathbf{C}_{10} = \{A_{22}, A_{23}\}$ and $|\mathbf{P}_{14}| = |\mathbf{P}_{23}| = 2$ indicating multiple inheritance.

At the gene level, we use $g_i \in A_j$ to denote that gene i is annotated by A_j . We further use $g_i \triangleleft A_j$ to indicate that A_j is the most specific GO term that annotates gene i , with the formal definition being $g_i \in A_j$ and $g_i \notin A_k$ for any $A_k \in \mathbf{C}_j$. We define $\mathbf{B}_i = \{A_j : g_i \triangleleft A_j\}$ to be the set of the most specific annotations of gene i . The *true-path* rule implies that \mathbf{B}_i contains all the annotation information about gene i . In Figure 1, we have $\mathbf{B}_a = \{A_{12}\}$ and $\mathbf{B}_k = \{A_2, A_{18}, A_{22}\}$. Let y_i ($i=1, \dots, I$) be the observed expression status of gene i , $y_i = 1$ if gene i is in S and $y_i = 0$ otherwise.

The binary y_i is assumed to follow a Bernoulli distribution, $y_i | p_i \sim \text{Bernoulli}(p_i)$, where p_i is the probability that gene i belongs to S . Using the idea that if A_j is associated with S , the genes annotated by A_j have a higher chance of being grouped in S , we construct the following logistic model,

$$\log\left(\frac{p_i}{1-p_i}\right) = b_0 + \sum_{j=1}^J I(A_j \in \mathbf{B}_i) \alpha_j + e_i. \quad (2)$$

We specify b_0 as a constant and set $b_0 = \log[p_0/(1-p_0)]$ with $p_0 = I_S/I$, where p_0 is the background probability that gene i is grouped in S by chance. The random error e_i is assumed to have a normal distribution with mean 0 and variance σ^2 , denoted by $e_i \stackrel{iid}{\sim} N(0, \sigma^2)$. Parameter α_j characterizes the relevance of GO term A_j to the set of identified genes S , where it modifies the odds of gene i being grouped in S by a factor of $\exp(\alpha_j)$ if $g_i \triangleleft A_j$. Thus, positive (negative) values of α_j indicate over(under)-representation. Based on the *true-path* rule, we only include α_j 's from \mathbf{B}_i , the most specific annotations, in model (2) to avoid repeated use of information. The α_j 's from less-specific annotations are assumed to affect the odds indirectly via a hierarchical prior on $\boldsymbol{\alpha} = \{\alpha_j, j=1, \dots, J\}$, constructed according to the dependence structure in the GO DAG. We set $\alpha_1 = 0$ for the root node. Then the prior of α_j ($j=2, \dots, J$) is specified conditionally given the relevance parameters of its parent nodes,

denoted by $\boldsymbol{\alpha}_{P_j} = \{\alpha_k : A_k \in \mathbf{P}_j\}$. Specifically,

$$\alpha_j | \boldsymbol{\alpha}_{P_j}, \delta^2 \sim \sum_{k: A_k \in \mathbf{P}_j} \frac{1}{|\mathbf{P}_j|} N(\alpha_k, \delta^2). \quad (3)$$

Prior (3) assumes that α_j arises from a mixture distribution of $|\mathbf{P}_j|$ components, each component being a normal distribution centered at the relevance parameter of one of its parents. The equal mixing probability $1/|\mathbf{P}_j|$ in (3) assumes a priori that each parent has equal influence over α_j . Parameter δ^2 characterizes the variability among the children nodes. The joint prior of $\boldsymbol{\alpha}$ is obtained by the product of (3) over $j=2, \dots, J$. This prior provides a mechanism to share information among the GO terms based on the DAG structure. It also naturally accommodates multiple inheritance.

We assign an inverse gamma prior, $\text{IG}(a_\delta, b_\delta)$, on δ^2 , which has been used extensively in Bayesian models (Gelman *et al.*, 2003). Similarly, an $\text{IG}(a_\sigma, b_\sigma)$ prior is assumed for σ^2 . We infer the relevance of GO term A_j based on the posterior distribution of α_j , denoted by $[\alpha_j | Y]$. Here, $Y = \{y_i, i=1, \dots, I\}$ is the collection of observations. Specifically, we use $r_j \equiv P(\alpha_j > 0 | Y)$ (denoted as the *B-score*) to measure the relevance of a GO term to S . It is the posterior probability of A_j being positively associated with S . Making inferences based on posterior probabilities is a common practice in Bayesian analysis of microarray data (Newton *et al.* 2004; Do *et al.*, 2005; Cao *et al.*, 2009). Markov Chain Monte Carlo (MCMC) sampling algorithm is employed to simulate random samples from the joint posterior distribution. We implement the adaptive-rejection sampling method to take advantage of the log-concave property of the full conditional distributions (Gilks and Wild, 1992). The computation can be performed efficiently.

2.2 Demonstration of GO-Bayes with the artificial DAG

GO-Bayes is applied to the artificial DAG (Fig. 1) to detect overrepresentation of the terms. We present the *B-score*, r_j , below the hypergeometric *P*-value. In the following discussion, we use association to refer to positive association between a GO term and S .

By incorporating the dependence structure of the GO DAG, GO-Bayes shows distinctive advantages over the hypergeometric test. First, GO-Bayes is capable of distinguishing terms with the same (I_{A_j}, n_{A_j}) . For example, GO-Bayes produces $r_{14} = 0.958$ and $r_{24} = 0.772$, for A_{14} and A_{24} , respectively, indicating that A_{14} is more likely to be associated with S based on the stronger evidence of overrepresentation of its neighboring nodes. Second, the *B-score* for all nodes has a range of 0–1 regardless of I_A . Thus, GO-Bayes can identify more specific GO terms as long as their neighboring nodes are consistently overrepresented. For example, the top two terms selected by GO-Bayes are A_{12} and A_{14} . Third, GO-Bayes also highlights underrepresentations as well as overrepresentations. In Figure 1, all the terms with $n_A = 0$ have the same *P*-value of 1.0. Under all the branches of A_2 , A_3 and A_4 , there are terms with a *P*-value of 1.0. GO-Bayes suggests that it is the terms under A_3 , whose r_j 's are close to 0, that are most underrepresented in S .

We have also compared the GO-Bayes *B-score* with the *elim* *P*-value (Alexa *et al.*, 2006) and the parent-child (union) *P*-value (Grossmann *et al.*, 2007) based on the artificial DAG. Due to the space limit, the results are presented in Supplementary Figure 3. In general, both the *elim* method and the parent-child method address the ‘dependency problem’ caused by overlapping annotations between parent-child GO terms. The *elim* method removes (or downweights) all genes annotated to a significantly enriched node from all its ancestors. By doing this, the method tends to identify strongly overrepresented GO terms that remain significant even after discounting evidence from their offsprings. For illustrative purpose, we set the *P*-value cutoff at 0.05 for the *elim* method. Thus, A_{12} is considered significantly overrepresented and genes (a, b, c) are removed from its ancestor terms A_5 and A_2 . Based on the *elim* *P*-value, the term A_{12} becomes the most significant, where A_2 ranks second and A_5 ranks sixth. The parent-child method computes the significance of a node conditional on the significance of its parents. It implements the idea by computing a hypergeometric *P*-value for each GO term in the context of its parent (treating the genes annotated by the parent term as the full

gene list). This approach tends to identify GO terms that show stronger overrepresentation compared with their parents. Take A_7 and A_{11} , for example, which have the same (I_A, n_A) . Based on the parent-child P -value, A_{11} has a higher rank (second) because its overrepresentation (2 out of 3) is stronger than its parent A_4 (3 out of 11). In contrast, A_7 has a lower rank (10th) because its overrepresentation (2 out of 3) is weaker compared with its parent A_2 (8 out of 11). The GO-Bayes method accounts for the parent-child relationship through the hierarchical prior. The strategy of borrowing information from neighboring terms allows GO-Bayes to detect moderate but consistent signals from closely related GO terms. As a result, A_7 has a higher rank than A_{11} ($r_7=0.802$ and $r_{11}=0.397$) because the overrepresentation in A_7 is corroborated by related terms in its neighborhood. With the biological truth unknown, there is no gold standard to compare the methods in real studies (Grossmann *et al.*, 2007). The results based on the artificial DAG, however, help to illustrate the distinctive characteristics of each method.

3 APPLICATION

3.1 Dataset

In order to test the utility of the GO-Bayes approach, we selected a gene expression microarray dataset in which a B-cell lymphoma cell line (Ramos) was stimulated either through the B-cell antigen receptor (BCR), CD40 or a combination of the two (Basso *et al.*, 2005). Specifically, we used an Affymetrix gene expression dataset selected from the GSE2350 series (GSM44051 to GSM44074) downloaded from the NCBI GEO database (<http://www.ncbi.nlm.nih.gov/projects/geo/>). The expression values of all six replicates under four experimental conditions were normalized by rows and a list of 3952 differentially expressed genes ($I=3952$) were selected using the Significance Analysis of Microarray (SAM) approach (Tusher *et al.*, 2001). The differentially expressed genes were subdivided into 20 groups based on their expression pattern by K -means clustering using Euclidean distance as the similarity metric.

Gene Cluster #7 (GC7) was chosen for detailed analysis because it promised to reveal some interesting biology about B-cell responses to receptor signaling. The 196 genes ($I_5=196$) present in GC7 showed a particularly interesting expression pattern: these genes were upregulated in response to BCR signaling alone; however, this upregulation was suppressed when CD40 signaling was included (Fig. 2). These treatment conditions mimic important biological responses of immature B cells (Hsueh and Scheuermann, 2000). Immature B cells must learn to distinguish between signals delivered by authentic pathogen-derived antigens and signals delivered by self-antigens. In the former case, B cells need to respond by productive proliferation and differentiation into immune effector cells. In the later case, B-cell responses need to be suppressed either through the induction of a state of unresponsiveness or apoptotic cell death. The two-signal hypothesis is one mechanism proposed to elicit either responsiveness or non-responsiveness to antigen exposure, which states that B cells receiving only one signal, through the BCR, will proliferate and die, but B cells receiving two signals both through the BCR and a co-stimulatory receptor-like CD40 will proliferate and survive due to suppression of the cell death response. Thus, gene present in GC7 are those genes whose expression is suppressed with the addition of CD40 signaling and could thus be involved either in the cell death response or in the induction of unresponsiveness.

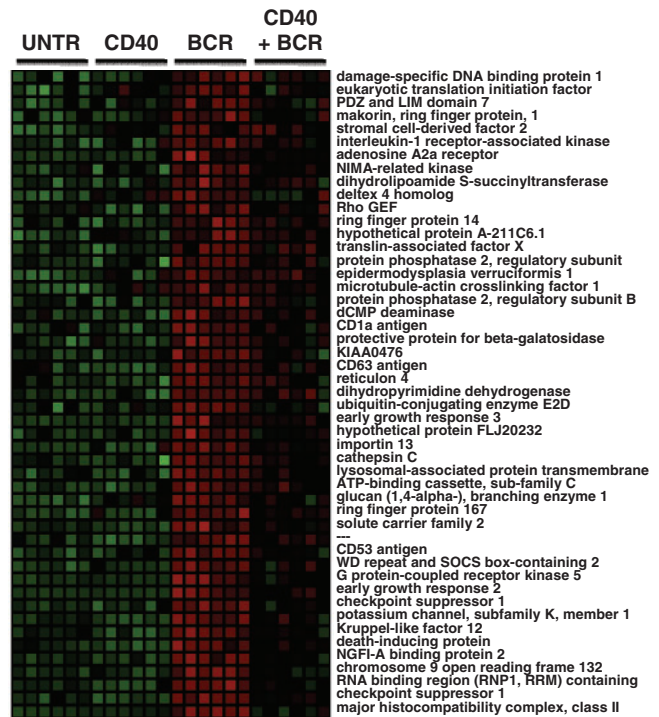


Fig. 2. Gene expression pattern of Gene Cluster #7. A heat map of normalized expression values is shown in which green represents relatively low expression and red represents relatively high expression. Each column represents data from a single microarray. Six replicates from each of four experiment conditions were performed—untreated (UNTR), stimulation through the CD40 receptor (CD40), stimulation through the BCR and the combined stimulation. The expression pattern for a subset of 196 genes in GC#7 is shown.

3.2 Result

For the full list of 3952 differentially expressed genes ($I=3952$), 6768 GO terms ($J=6768$) have been used to annotate their functions. Four groups of GO terms are found in the top 20 list of most significant terms based on the GO-Bayes approach (Table 1). The GO term with the highest B -score (0.9004) is ‘transferase activity, transferring hexosyl groups’ (GO:0016758). Nine of the top 20 GO terms based on the B -score are related to this group. Five of the top 20 GO terms are related to ‘G-protein signaling, coupled to cAMP nucleotide second messenger’ (GO:0007188); six are related to ‘lysosome’ (GO:0005764); and two are in the group of ‘negative regulation of transcription, DNA-dependent’ (GO:0045892). In the case of the ‘transferase activity, transferring hexosyl groups’, all nine related terms (GO:0016758, GO:0016757, GO:0000030, GO:0003844, GO:0008375, GO:0015020, GO:0042328, GO:0004703, GO:0004674) can be found near GO:0016758 in the GO hierarchy (see Supplementary Fig. 4). Even though some of these terms are poorly represented in this dataset (e.g. GO:0003844 represented by only a single gene in the dataset), all of these terms have relatively high B -scores due to the overrepresentation of other terms that are close relatives in the hierarchy.

In some cases, the association of the group of GO terms highlighted by the GO-Bayes approach matches our expectations.

Table 1. The top 20 lists GO terms associated with Gene Cluster #7 by CLASSIFI and GO-Bayes

| GO ID | I_A | n_A | P -value | Rank _P | B -score | Rank _B | Group | GO name |
|------------|-------|-------|------------|-------------------|------------|-------------------|-------|--|
| GO:0019933 | 21 | 6 | 4.00E-04 | 1 | 0.7652 | 149 | a | cAMP-mediated signaling |
| GO:0005773 | 91 | 13 | 4.61E-04 | 2 | 0.7544 | 187 | b | Vacuole |
| GO:0001726 | 23 | 6 | 6.85E-04 | 3 | 0.7724 | 118 | b | Ruffle |
| GO:0007188 | 16 | 5 | 7.95E-04 | 4 | 0.8688 | 7 | a | G-protein signaling, coupled to cAMP nucleotide second messenger |
| GO:0007190 | 10 | 4 | 9.73E-04 | 5 | 0.6928 | 508 | a | Activation of adenylate cyclase activity |
| GO:0045762 | 10 | 4 | 9.73E-04 | 5 | 0.6582 | 778 | a | Positive regulation of adenylate cyclase activity |
| GO:0031281 | 10 | 4 | 9.73E-04 | 5 | 0.4976 | 3148 | a | Positive regulation of cyclase activity |
| GO:0051349 | 10 | 4 | 9.73E-04 | 5 | 0.4712 | 3545 | a | Positive regulation of lyase activity |
| GO:0019935 | 25 | 6 | 1.11E-03 | 9 | 0.7568 | 177 | a | Cyclic-nucleotide-mediated signaling |
| GO:0007189 | 5 | 3 | 1.12E-03 | 10 | 0.8346 | 27 | a | G-protein signaling, adenylate cyclase activating pathway |
| GO:0016757 | 57 | 9 | 1.71E-03 | 11 | 0.8716 | 6 | c | Transferase activity, transferring glycosyl groups |
| GO:0000323 | 80 | 11 | 1.74E-03 | 12 | 0.8276 | 33 | b | Lytic vacuole |
| GO:0005764 | 80 | 11 | 1.74E-03 | 12 | 0.8762 | 5 | b | Lysosome |
| GO:0010324 | 69 | 10 | 1.87E-03 | 14 | 0.5084 | 2957 | b | Membrane invagination |
| GO:0006897 | 69 | 10 | 1.87E-03 | 14 | 0.5596 | 2015 | b | Endocytosis |
| GO:0007187 | 20 | 5 | 2.40E-03 | 16 | 0.8220 | 43 | a | G-protein signaling, coupled to cyclic nucleotide second messenger |
| GO:0006898 | 20 | 5 | 2.40E-03 | 16 | 0.5870 | 1601 | b | Receptor-mediated endocytosis |
| GO:0001609 | 2 | 2 | 2.45E-03 | 18 | 0.2706 | 6218 | a | Adenosine receptor activity, G-protein coupled |
| GO:0032230 | 2 | 2 | 2.45E-03 | 18 | 0.7562 | 178 | f | Positive regulation of synaptic transmission, GABAergic |
| GO:0048285 | 2 | 2 | 2.45E-03 | 18 | 0.5246 | 2625 | f | Organelle fission |
| GO:0007217 | 2 | 2 | 2.45E-03 | 18 | 0.8452 | 17 | a | Tachykinin signaling pathway |
| GO:0051319 | 2 | 2 | 2.45E-03 | 18 | 0.5124 | 2878 | f | G2-phase |
| GO:0015851 | 2 | 2 | 2.45E-03 | 18 | 0.3764 | 5103 | f | Nucleobase transport |
| GO:0016519 | 2 | 2 | 2.45E-03 | 18 | 0.1972 | 6600 | f | Gastric inhibitory peptide receptor activity |
| GO:0000085 | 2 | 2 | 2.45E-03 | 18 | 0.5650 | 1945 | f | G2-phase of mitotic cell cycle |
| GO:0016021 | 874 | 59 | 4.68E-03 | 31 | 0.8924 | 3 | a, b | Integral to membrane |
| GO:0005886 | 727 | 50 | 6.95E-03 | 36 | 0.8562 | 9 | a, b | Plasma membrane |
| GO:0004703 | 3 | 2 | 7.10E-03 | 40 | 0.8536 | 15 | c | G-protein-coupled receptor kinase activity |
| GO:0000030 | 4 | 2 | 1.37E-02 | 61 | 0.8556 | 11 | c | Mannosyltransferase activity |
| GO:0016758 | 41 | 6 | 1.45E-02 | 70 | 0.9004 | 1 | c | Transferase activity, transferring hexosyl groups |
| GO:0015020 | 5 | 2 | 2.22E-02 | 92 | 0.8538 | 14 | c | glucuronosyltransferase activity |
| GO:0045892 | 89 | 9 | 3.09E-02 | 104 | 0.8996 | 2 | d | Negative regulation of transcription, DNA dependent |
| GO:0015075 | 135 | 12 | 3.41E-02 | 115 | 0.8568 | 8 | b | Ion transmembrane transporter activity |
| GO:0003844 | 1 | 1 | 4.96E-02 | 281 | 0.8506 | 16 | c | 1,4-Alpha-glucan branching enzyme activity |
| GO:0022891 | 160 | 13 | 5.22E-02 | 289 | 0.8562 | 10 | b | Substrate-specific transmembrane transporter activity |
| GO:0042328 | 2 | 1 | 9.67E-02 | 471 | 0.8414 | 20 | c | Heparan sulfate <i>N</i> -acetylglucosaminyltransferase activity |
| GO:0000122 | 66 | 6 | 1.07E-01 | 491 | 0.8556 | 12 | d | Negative regulation of transcription from RNA polymerase II promoter |
| GO:0004674 | 200 | 14 | 1.18E-01 | 512 | 0.8550 | 13 | c | Protein serine/threonine kinase activity |
| GO:0008375 | 14 | 2 | 1.51E-01 | 620 | 0.8770 | 4 | c | Acetylglucosaminyltransferase activity |
| GO:0007186 | 101 | 7 | 2.33E-01 | 814 | 0.8430 | 19 | a | G-protein-coupled receptor protein signaling pathway |
| GO:0042629 | 2 | 0 | 1.0 | NA | 0.8432 | 18 | b | Mast cell granule |

The columns I_A and n_A represent the cardinality of the relevant GO term and the number of genes in the GO term that also appear in Gene Cluster #7; P -value and Rank_P represent the hypergeometric P -value computed by CLASSIFI, and the rank of the GO terms based on the P -value; and B -score and Rank_B represent the GO-Bayes measure and the corresponding rank of the GO terms. Under the column 'Group', 'a' represents GO terms closely related to GO:0007188 (G-protein signaling, coupled to cAMP nucleotide second messenger) in the GO hierarchy; 'b' represents GO terms closely related to GO:0005764 (lysosome); 'c' represents GO terms closely related to GO:0016757 (transferase activity, transferring glycosyl groups); 'd' represents GO terms closely related to GO:0045892 (negative regulation of transcription, DNA-dependent); and 'f' represents GO terms unrelated to the above four groups. NA, not applicable.

It is well known that activation of B cells through the BCR induces endocytosis and the fusion of endocytic vesicles with lysosomes as a mechanism to capture antigen for presentation to T cells in order to stimulate the helper immune response (Lee *et al.*, 2006). Thus, the presence of lysosome-related GO terms in the cluster of genes upregulated in response to BCR stimulation might be expected. In other cases, the association is not immediately expected, but subsequent investigations revealed that the association is supported by previous experiment data. For example, the majority

of genes giving rise to the 'negative regulation of transcription, DNA-dependent' association, including ID3 (Pan *et al.*, 1999), CTCF (Qi *et al.*, 2003), SMAD3 (Ramesh *et al.*, 2009), KLF12 (Roth *et al.*, 2000), E2F6 (Xu *et al.*, 2007), FosB (Yin *et al.*, 2007) and PA2G4 (Zhang *et al.*, 2008), have been found to be upregulated in response to BCR stimulation in B cells or somehow involved in B-cell signaling responses. In the case of CTCF, Qi *et al.* (2003) showed that this upregulation is associated with the induction of apoptosis in B cells and can be suppressed by co-stimulation through

CD40 in agreement with the findings reported here. In still other cases, no direct corroborative evidence could be found (e.g. for ‘transferase activity, transferring hexosyl groups’). Thus, this finding serves as a hypothesis for future testing.

In order to compare the results of the GO-Bayes approach with the standard ORA based on the hypergeometric test, we processed the gene set in GC7 with the CLASSIFI algorithm (Lee *et al.*, 2006) and selected the top 20 GO terms. The top 20 lists by CLASSIFI and GO-Bayes, respectively, are presented in Table 1. While four GO terms were ranked in the top 20 by both methods, the remaining top 20 terms from each method were distinct. Based on this comparison, several distinctions between the two approaches can be made.

First, three of the four GO term groups are found in both top 20 term lists. Two of the groups, G protein signaling and lysosome, have multiple terms in both top 20 lists. One group, transferase activity, is only represented once in the CLASSIFI list, but multiple times in the GO-Bayes list. Given that little could be found in the literature about these genes in B cell biology, this association would likely be ignored from the hypergeometric analysis. One group, the negative regulation of transcription group was not found in the hypergeometric top 20, and yet is potentially the most interesting given the strong literature support described above.

Second, eight GO terms in GC7 with $(I_A, n_A) = (2, 2)$ have the same hypergeometric P -value of 0.0024, and all of them are in the hypergeometric top 20. The GO-Bayes measure suggests that they are very different in their association with the cluster. For example, GO:0007217 (tachykinin signaling pathway) has a B -score of 0.8452 and it ranks in the GO-Bayes top 20. In contrast, GO:0016519 (gastric inhibitory peptide receptor activity) has a B -score of 0.1972 and its GO-Bayes rank is 6600. To shed light on the difference in the B -scores between these two terms, we compare their regional DAGs, which are shown in Supplementary Figures 5 and 6. The three direct ancestors of GO:0007217 have a stronger association with the cluster than those of GO:0016519. In addition, GO:0007217 has nine siblings, five of which have genes represented in the cluster ($n_A > 0$). By comparison, GO:0016519 has 13 siblings, of which only one sibling has genes represented in the cluster. The support from the related GO terms is substantially higher for GO:0007217 than for GO:0016519, and thus GO:0007217 is judged to be more likely associated with the cluster based on the GO-Bayes measure. While there is no evidence for the involvement of the gastric inhibitory peptide receptor activity in the regulation of B-cell function in the literature, tachykinin (also known as hemokinin-1) has been found to be secreted during the differentiation of B-cell precursors thereby regulating their own development (Milne *et al.*, 2004). Thus, GO:0007217 does appear to be biologically associated with GC7.

Among the top 20 GO terms identified by GO-Bayes, 8 GO terms have two genes or fewer belonging to GC7 (GO: 0007217, GO:0004703, GO:0000030, GO:0015020, GO:0003844, GO:42328, GO:0008375 and GO:0042629). Researchers may have concern over these findings, suspecting that they may be false positives. Based on the GO hierarchy, GO: 0007217 is a child of GO:0007186, which ranked seventh in the top 20 GO-Bayes list. The term is also found to be biologically associated with the experiment (Milne *et al.*, 2004). Six of the GO terms (GO:0004703, GO:0000030, GO:0015020, GO:0003844, GO:42328 and GO:0008375) are all related to ‘transferase activity’, and all of them are in a close neighborhood in the GO DAG as shown in Supplementary Figure 4. As for GO:0042629, the

fact that none of the two annotated genes were found in GC7 requires in-depth discussion on the term. Note that GO:0042629 is a direct child of ‘lysosome’ (GO:0005764), which ranked fifth in the top 20 GO-Bayes list. The two genes annotated with GO:0042629 are IL8 receptor beta (IL8RB) and serglycin (SRGN). IL8RB has been found to be involved in B-cell chemotaxis across the blood-brain barrier (Alter *et al.*, 2003). Mice deficient for IL8RB show lymphadenopathy due to a specific expansion of the B-cell compartment (Cacalano *et al.*, 1994). Importantly, IL8RB was found to be dramatically upregulated (~9-fold) in response to BCR stimulation of primary splenic B cells by microarray analysis (Sato *et al.*, 2005) (see IL8RB expression profile in <http://www.ncbi.nlm.nih.gov/sites/GDSbrowser?acc=GDS1467>). Although less is known about the function of SRGN in B cells, it was also found to be upregulated following BCR stimulation of B cells from TAK1-deficient mice in the same microarray experiment. Taken together, these findings suggest that the identification of GO:0042629 by GO-Bayes is significant and represents an example where the GO-Bayes approach was able to overcome false negative results generated by deficiencies in microarray data generation or processing. Thus, GO-Bayes will only identify a GO term with relatively few annotated genes belonging to S when there is strong evidence of overrepresentation from its neighboring GO terms.

4 DISCUSSION

In this article, we proposed GO-Bayes, a Bayesian approach for GO-based OR. The model has a relatively simple format, with the first level as a logistic regression model and straightforward prior specification. The key innovation is that it can incorporate the dependence structure of the GO DAG on a global scale. The resulting measure on the association between GO terms and selected genes borrows information across related GO terms to strengthen the detection of overrepresentation signals. We have given detailed comparison between the GO-Bayes approach and the hypergeometric test, which is widely used in ORA. Our analysis using an artificial dataset and a real microarray dataset suggests that the GO-Bayes approach can produce more biologically meaningful results than the hypergeometric test.

Relying on individual GO terms, the hypergeometric P -value has a closed form, which only requires the information on I , I_S and (I_A, n_A) for each GO term. GO-Bayes, on the other hand, does not have a closed form and it needs the information on the structure of the GO DAG. We have developed a program in C to implement the GO-Bayes approach. For the B-cell lymphoma cell line example in the application section, the B -score calculations were completed in ~10 min on a MAC (OS X 10.4.11, 1.83 GHz Intel Core Duo processor) computer. The program is available upon request from S.Z.

The GO-Bayes measure, $r_j \equiv P(\alpha_j > 0 | Y)$, is the posterior probability of a GO term being positively associated with S . As a reviewer pointed out, r_j behaves like a frequentist P -value under the null hypothesis (Bochkina and Richardson, 2007), and a frequentist-type estimator of false discovery rate (Storey, 2002) can be employed to assess the statistical significance.

Currently, the proposed method is based on binary outcomes (membership of genes in S). We are working on generalizing the GO-Bayes approach to cases with ordinal or continuous outcomes. Although we have focused on the use of this approach

for the interpretation of gene expression microarray data and GO annotation, the general strategy can be applied to any circumstance in which groups of entities are annotated with terms derived from any ontology hierarchy or similar dependency structure.

ACKNOWLEDGEMENTS

The authors thank the three reviewers for their constructive comments and suggestions.

Funding: U.S. National Institutes of Health (N01AI40076 and UL1 RR024982, in part).

Conflict of Interest: none declared.

REFERENCES

- Alexa, A. *et al.* (2006) Improved scoring of functional groups from gene expression data by decorrelating GO graph structure. *Bioinformatics*, **22**, 1600–1607.
- Alter, A. *et al.* (2003) Determinants of human B cell migration across brain endothelial cells. *J. Immunol.*, **170**, 4497–4505.
- Al-Shahrour, F. *et al.* (2004) FatiGO: a web tool for finding significant associations of Gene Ontology terms with groups of genes. *Bioinformatics*, **20**, 578–580.
- Ashburner, M. *et al.* (2000) Gene ontology: tool for the unification of biology. The Gene Ontology Consortium. *Nat. Genet.*, **25**, 25–29.
- Basso, K. *et al.* (2005) Reverse engineering of regulatory networks in human B cells. *Nat. Genet.*, **17**, 182–190.
- Beissbarth, T. and Speed, T.P. (2004) GStat: find statistically overrepresented gene ontologies within a group of genes. *Bioinformatics*, **20**, 1464–1465.
- Bochkina, N. and Richardson, S. (2007) Tail posterior probability for inference in pairwise and multiclass gene expression data. *Biometrics*, **63**, 1117–1125.
- Cacalano, G. *et al.* (1994) Neutrophil and B cell expansion in mice that lack the murine IL-8 receptor homolog. *Science*, **265**, 682–684.
- Cao, J. *et al.* (2009) Bayesian optimal discovery procedure for simultaneous significance testing. *BMC Bioinformatics*, **10**, 5.
- Cho, R.J. *et al.* (2001) Transcriptional regulation and function during the human cell cycle. *Nat. Genet.*, **27**, 48–54.
- Do, K. *et al.* (2005) A Bayesian mixture model for differential gene expression. *Appl. Stat.*, **54**, 627–644.
- Drăghici, S. *et al.* (2003) Global functional profiling of gene expression. *Genomics*, **81**, 98–104.
- Grossmann, S. *et al.* (2007) Improved detection of overrepresentation of Gene-Ontology annotations with parent-child analysis. *Bioinformatics*, **23**, 3024–3031.
- Gelman, A. *et al.* (2003) *Bayesian Data Analysis*. CRC Press, London.
- Gilks, W. and Wild, P. (1992) Adaptive rejection sampling for Gibbs sampling. *Appl. Stat.*, **41**, 337–348.
- Hsueh, R. and Scheuermann, R.H. (2000) Tyrosine kinase activation in the growth, differentiation and death responses initiated from the B cell antigen receptor. *Adv. Immunol.*, **75**, 283–316.
- Khatri, P. *et al.* (2002) Profiling gene expression using Onto-Express. *Genomics*, **79**, 266–270.
- Khatri, P. and Drăghici, S. (2005) Ontological analysis of gene expression data: current tools, limitations, and open problems. *Bioinformatics*, **21**, 3587–3595.
- Lee, H.K. *et al.* (2005) ErmineJ: tool for functional analysis of gene expression data sets. *BMC Bioinformatics*, **6**, 269.
- Lee, J.A. *et al.* (2006) Components of the antigen processing and presentation pathway revealed by gene expression microarray analysis following B cell antigen receptor (BCR) stimulation. *BMC Bioinformatics*, **7**, 237.
- Luo, F. *et al.* (2007) Modular organization of protein Interaction networks. *Bioinformatics*, **23**, 207–214.
- Lewin, A.M. and Grieve, I.C. (2006) Grouping Gene Ontology terms to improve the assessment of gene set enrichment in microarray data. *BMC Bioinformatics*, **7**, 426.
- Milne, C.D. *et al.* (2004) Mechanisms of selection mediated by interleukin-7, the preBCR, and hemokinin-1 during B-cell development. *Immunol. Rev.*, **197**, 75–88.
- Newton, M.A. *et al.* (2004) Detecting differential gene expression with a semiparametric hierarchical mixture method. *Biostatistics*, **4**, 155–176.
- Pan, L. *et al.* (1999) Impaired Immune Responses and B-Cell Proliferation in Mice Lacking the Id3 Gene. *Mol. Cell. Biol.*, **19**, 5969–5980.
- Qi, C. *et al.* (2003) CTCF functions as a critical regulator of cell-cycle arrest and death after ligation of the B cell receptor on immature B cells. *Proc. Natl Acad. Sci. USA*, **100**, 633–638.
- Ramesh, S. *et al.* (2009) Transforming growth factor β (TGF β)-induced apoptosis. *Cell Cycle*, **8**, 11–17.
- Roth, C. *et al.* (2000) Genomic structure and DNA binding properties of the human zinc finger transcriptional repressor AP-2rep (KLF12). *Genomics*, **63**, 384–390.
- Sato, S. *et al.* (2005) Essential function for the kinase TAK1 in innate and adaptive immune responses. *Nat. Immunol.*, **6**, 1087–1095.
- Storey, J.D. (2002). A direct approach to false discovery rate. *J. R. Stat. Soc. Ser. B*, **64**, 479–498.
- Tusher, V.G. *et al.* (2001) Significance analysis of microarrays applied to transcriptional responses to ionizing radiation. *Proc. Natl Acad. Sci. USA*, **98**, 5116–5121.
- Xu, X. *et al.* (2000) A comprehensive ChIP-chip analysis of E2F1, E2F4, and E2F6 in normal and tumor cells reveals interchangeable roles of E2F family members. *Genome Res.*, **17**, 1550–1561.
- Yin, Q. *et al.* (2007) B-cell receptor activation induces BIC/miR-155 expression through a conserved AP-1 element. *J. Biol. Chem.*, **283**, 2654–2662.
- Zhang, Y. *et al.* (2008) Alterations in cell growth and signaling in Erbb3 binding protein-1 (Ebp1) deficient mice. *BMC Cell Biol.*, **9**, 69.